



CSCI-GA.2250-001

Operating Systems

Networking

Hubertus Franke
frankeh@cs.nyu.edu



TCP/IP protocol family

- IP : Internet Protocol
 - UDP : User Datagram Protocol
 - RTP, traceroute
 - TCP : Transmission Control Protocol
 - HTTP, FTP, ssh

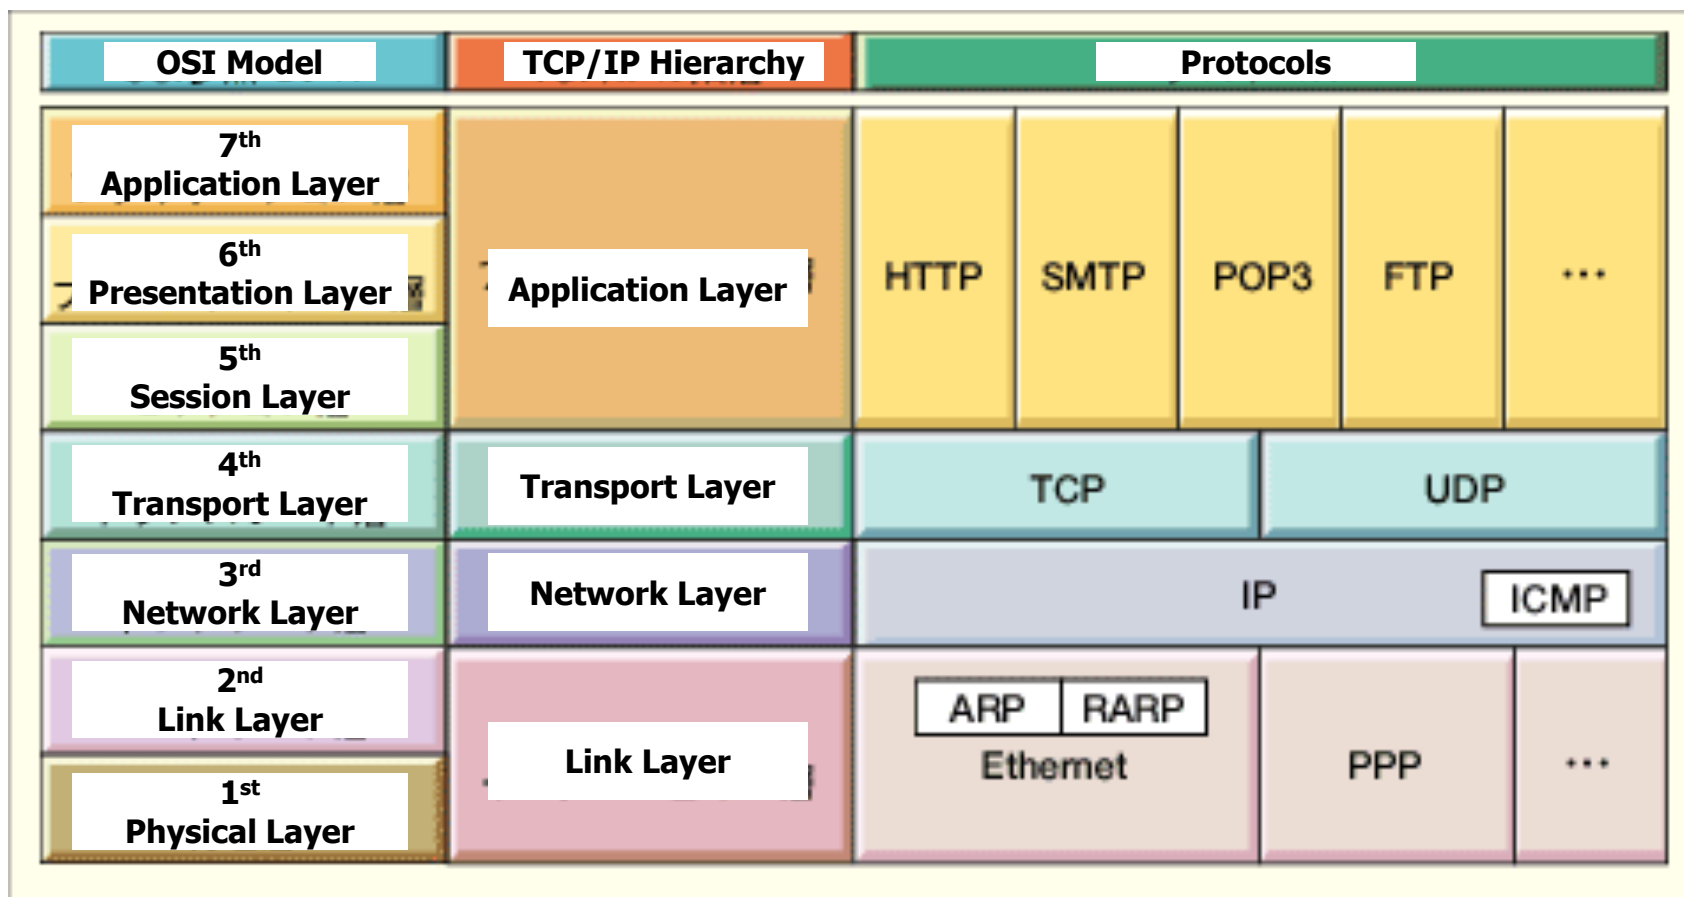
What is an internet?

- A set of *interconnected networks*
- The Internet is the most famous example
- Networks can be completely different
 - Ethernet, ATM, modem, ...
 - (TCP/)IP is what links them

What is an internet? (cont)

- *Routers* (nodes) are devices on multiple networks that pass traffic between them
- Individual networks pass traffic from one router or endpoint to another
- TCP/IP hides the details as much as possible

OSI and Protocol Stack



Link Layer : includes device driver and network interface card

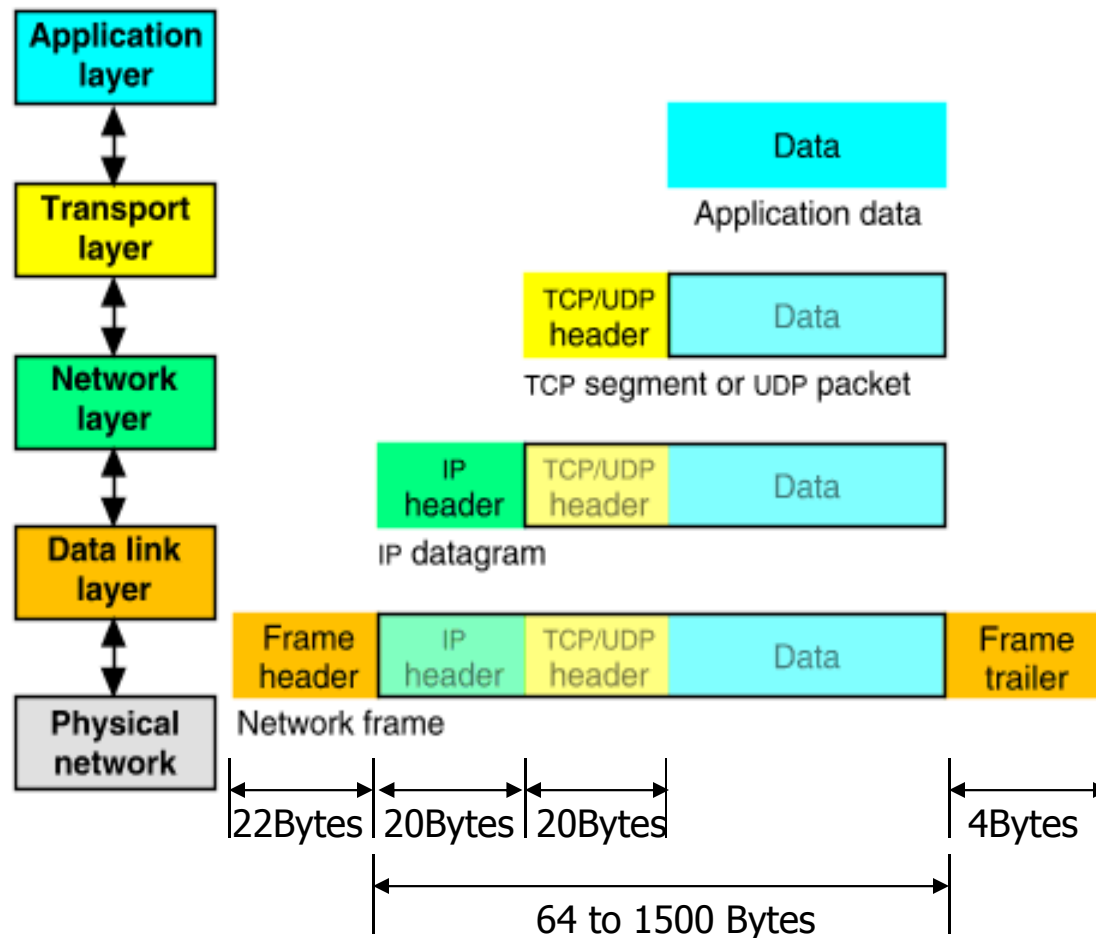
Network Layer : handles the movement of packets, i.e. Routing

Transport Layer : provides a reliable flow of data between two hosts

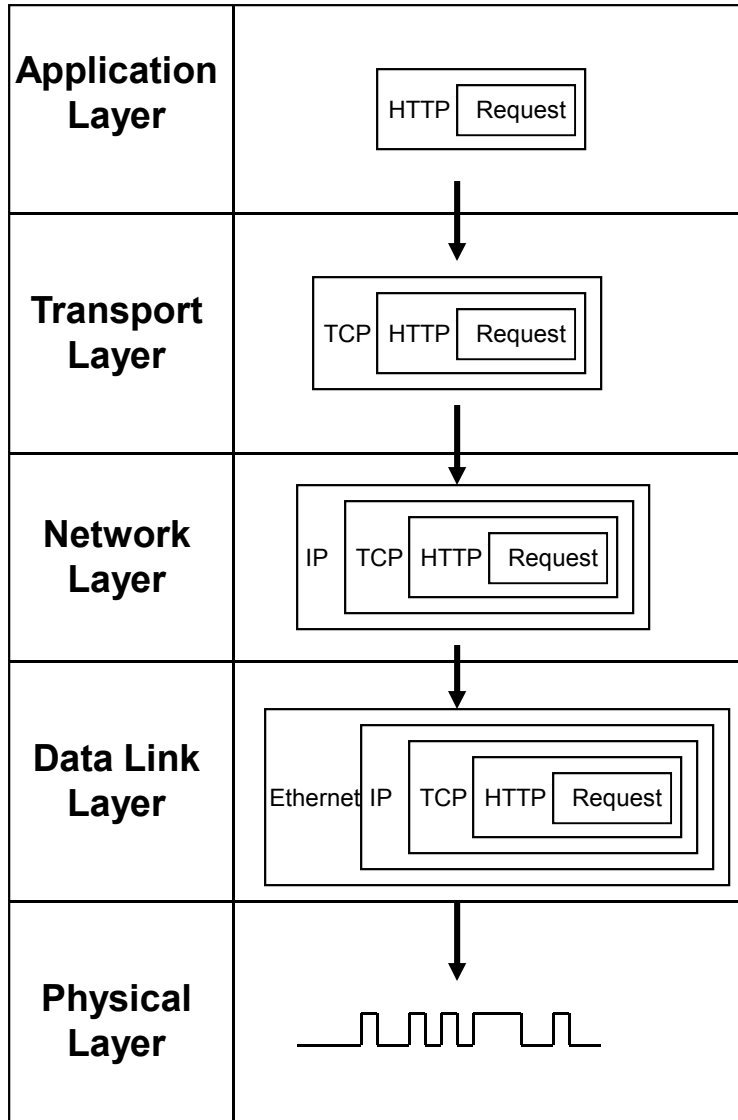
Application Layer : handles the details of the particular application

Packet Encapsulation

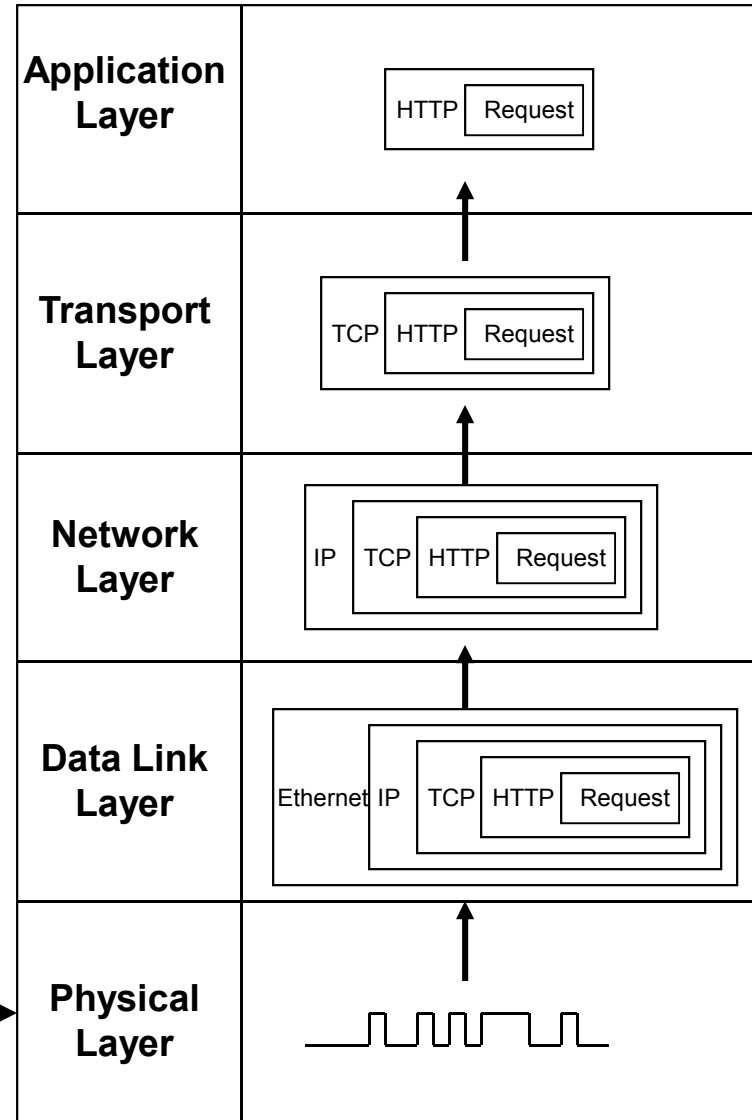
- The data is sent down the protocol stack
- Each layer adds to the data by prepending headers



Sender



Receiver



IP

- Responsible for end to end transmission
- Sends data in individual packets
- Maximum size of packet is determined by the networks
 - Fragmented if too large
- Unreliable
 - Packets might be lost, corrupted, duplicated, delivered out of order

IP addresses

- 4 bytes
 - e.g. 163.1.125.98
 - Each device normally gets one (or more)
 - In theory there are about 4 billion available
- But...


Routing

- How does a device know where to send a packet?
 - All devices need to know what IP addresses are on directly attached networks
 - If the destination is on a local network, send it directly there

Routing (cont)

- If the destination address isn't local
 - Most non-router devices just send everything to a single local router
 - Routers need to know which network corresponds to each possible IP address

Routing



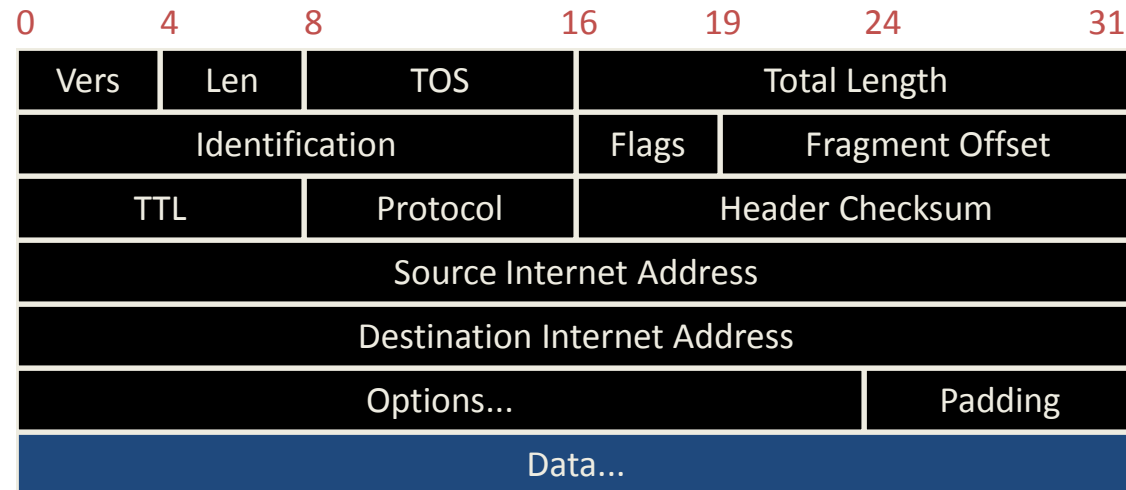
Allocation of addresses

- Controlled centrally by ICANN
 - Fairly strict rules on further delegation to avoid wastage
 - Have to demonstrate actual need for them
- Organizations that got in early have bigger allocations than they really need

IP packets

- Source and destination addresses
- Protocol number
 - 1 = ICMP, 6 = TCP, 17 = UDP
- Various options
 - e.g. to control fragmentation
- Time to live (TTL)
 - Prevent routing loops

IP Datagram



Field	Purpose
-------	---------


Vers	IP version number
Len	Length of IP header (4 octet units)
TOS	Type of Service
T. Length	Length of entire datagram (octets)
Ident.	IP datagram ID (for frag/reassembly)
Flags	Don't/More fragments
Frag Off	Fragment Offset

Field	Purpose
-------	---------

TTL	Time To Live - Max # of hops
Protocol	Higher level protocol (1=ICMP, 6=TCP, 17=UDP)
Checksum	Checksum for the IP header
Source IA	Originator's Internet Address
Dest. IA	Final Destination Internet Address
Options	Source route, time stamp, etc.
Data...	Higher level protocol data

We only looked at the IP addresses, TTL and protocol #

IP Routing




- Routing Table
 - Destination IP address
 - IP address of a next-hop router
 - Flags
 - Network interface specification

UDP

- Thin layer on top of IP
- Adds packet length + checksum
 - Guard against corrupted packets
- Also source and destination *ports*
 - Ports are used to associate a packet with a specific application at each end
- Still unreliable:
 - Duplication, loss, out-of-orderness possible

UDP datagram



Field	Purpose
Source Port	16-bit port number identifying originating application
Destination Port	16-bit port number identifying destination application
Length	Length of UDP datagram (UDP header + data)
Checksum	Checksum of IP pseudo header, UDP header, and data

Typical applications of UDP

- Where packet loss etc is better handled by the application than the network stack
 - Where the overhead of setting up a connection isn't wanted
-
- VOIP
 - NFS – Network File System
 - Most games

TCP

- Reliable, *full-duplex*, *connection-oriented*, *stream* delivery
 - Interface presented to the application doesn't require data in individual packets
 - Data is guaranteed to arrive, and in the correct order without duplications
 - Or the connection will be dropped
 - Imposes significant overheads

Applications of TCP

- Most things!
 - HTTP, FTP, ...
- Saves the application a lot of work, so used unless there's a good reason not to

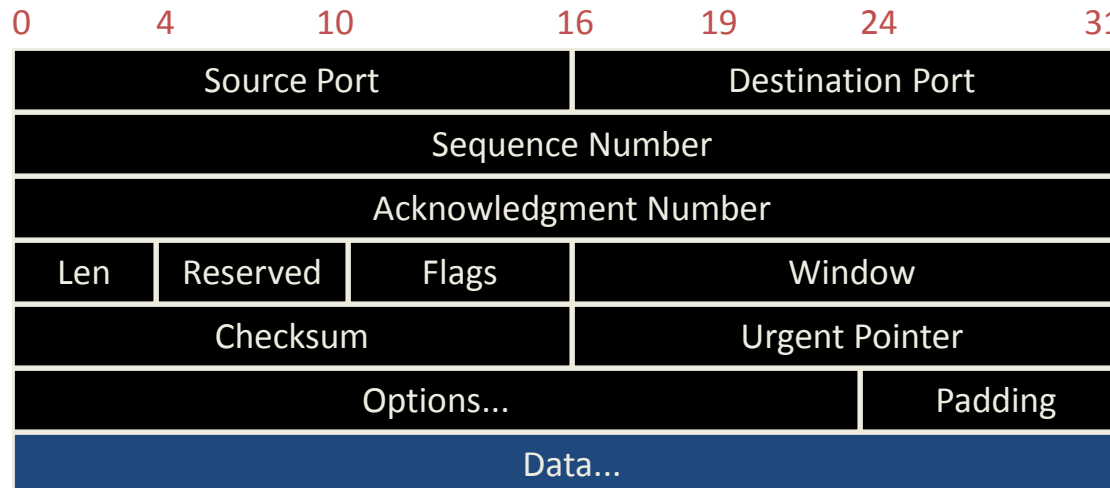
TCP implementation

- Connections are established using a *three-way handshake*
- Data is divided up into packets by the operating system
- Packets are numbered, and received packets are acknowledged
- Connections are explicitly closed
 - (or may abnormally terminate)

TCP Packets

- Source + destination ports
- Sequence number (used to order packets)
- Acknowledgement number (used to verify packets are received)

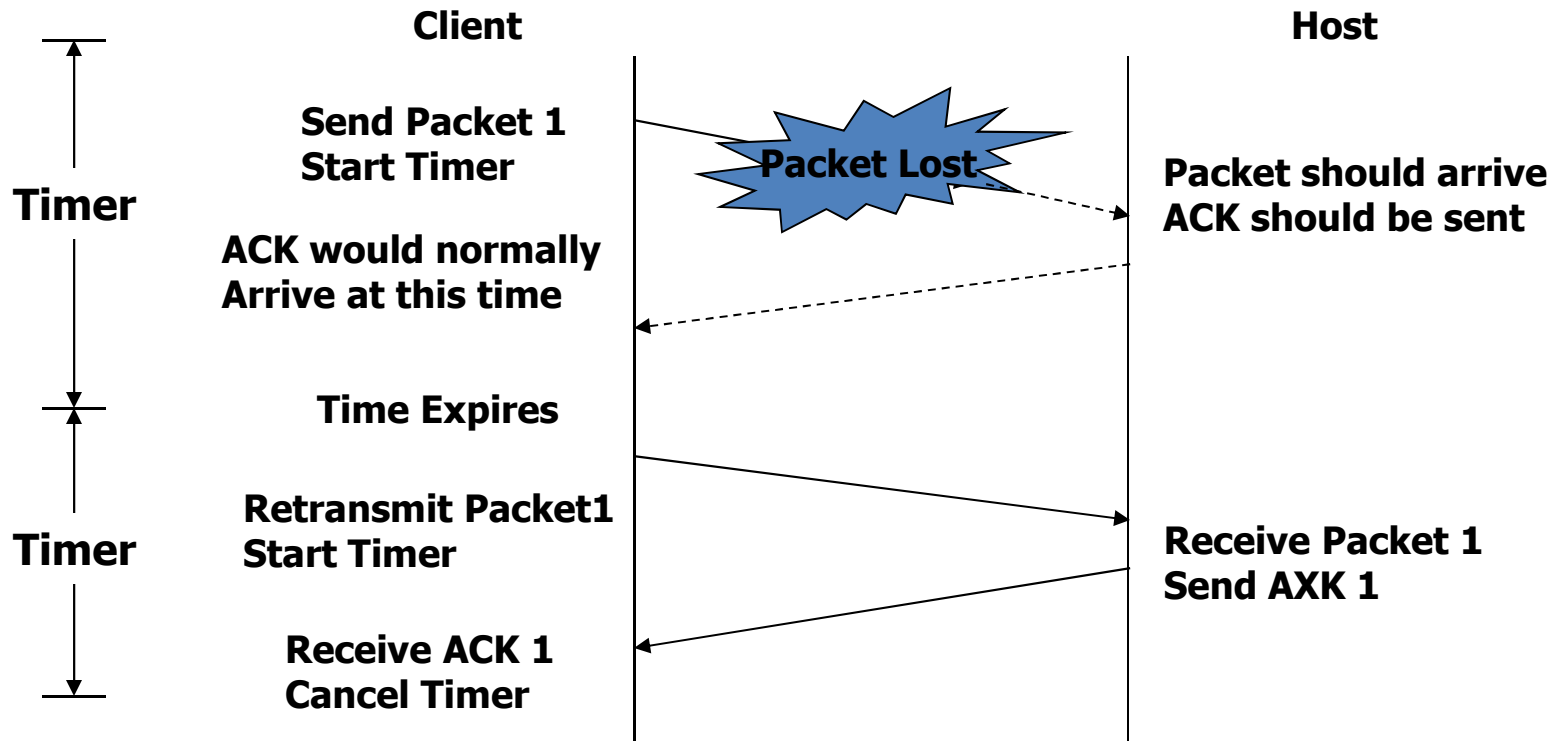
TCP Segment



Field	Purpose
Source Port	Identifies originating application
Destination Port	Identifies destination application
Sequence Number	Sequence number of first octet in the segment
Acknowledgment #	Sequence number of the next expected octet (if ACK flag set)
Len	Length of TCP header in 4 octet units
Flags	TCP flags: SYN, FIN, RST, PSH, ACK, URG
Window	Number of octets from ACK that sender will accept
Checksum	Checksum of IP pseudo-header + TCP header + data
Urgent Pointer	Pointer to end of "urgent data"
Options	Special TCP options such as MSS and Window Scale

You just need to know port numbers, seq and ack are added

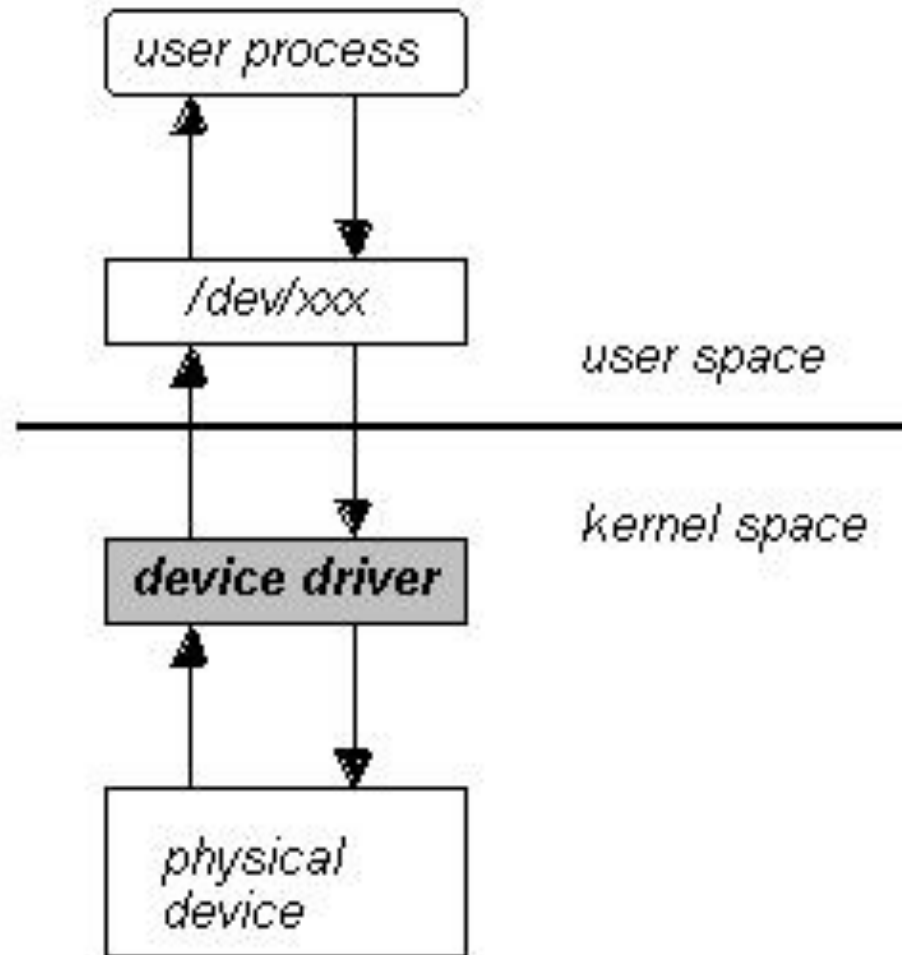
TCP : Data transfer



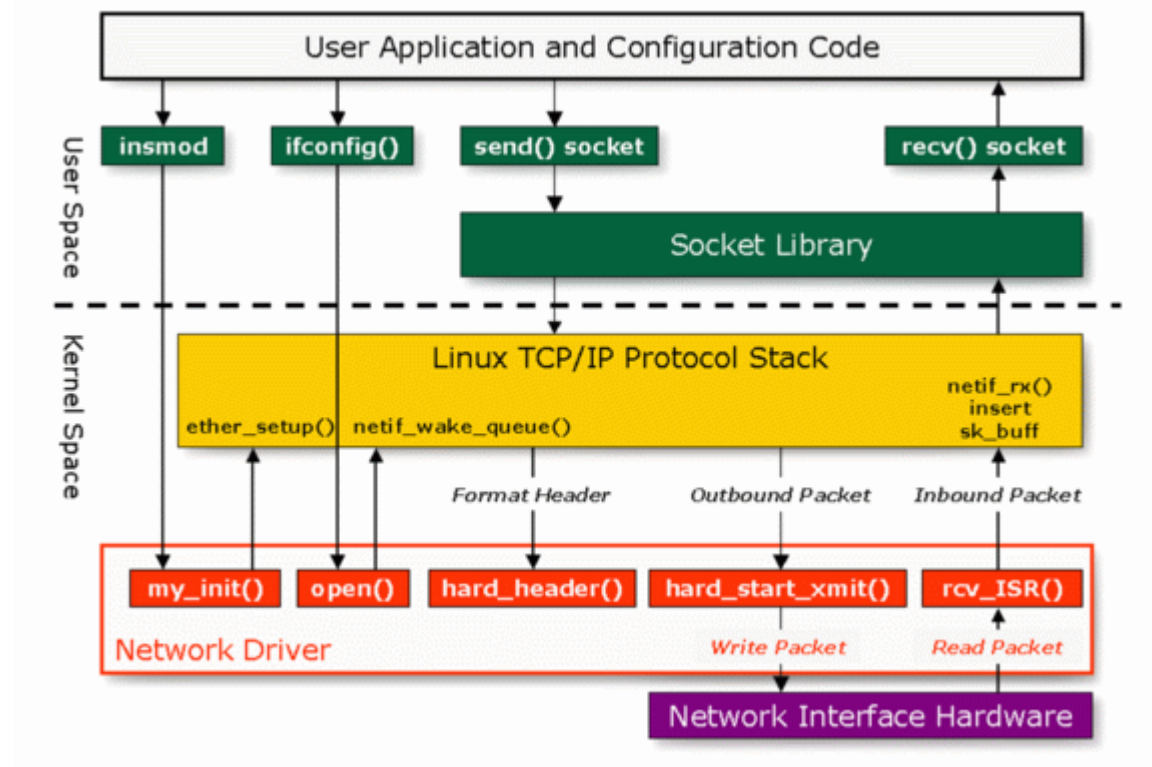
IPv6

- 128 bit addresses
 - Make it feasible to be very wasteful with address allocations
- Lots of other new features
 - Built-in autoconfiguration, security options, ...
- Not really in production use yet

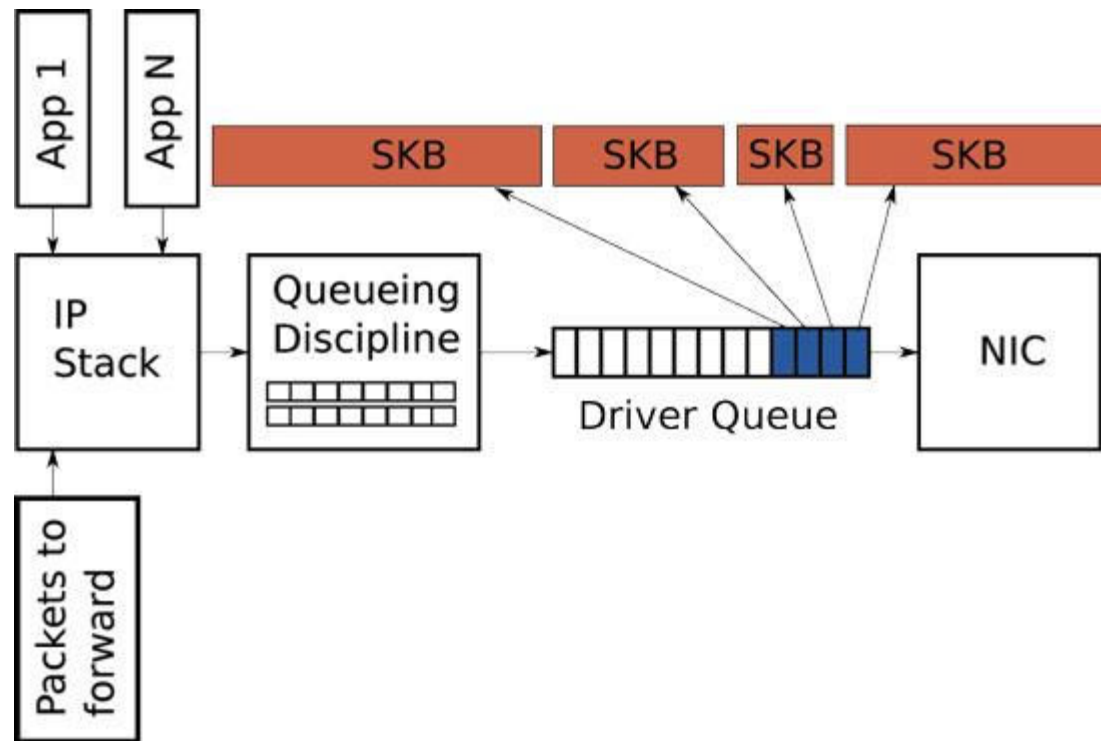
General Structure



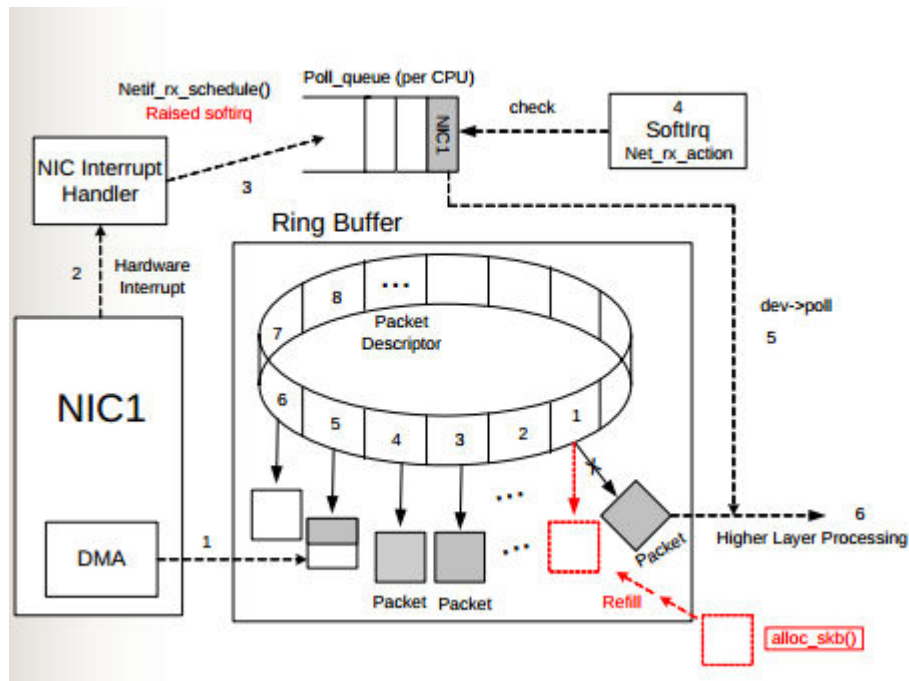
Anatomy of Network Stack



Interaction IP Stack / NIC

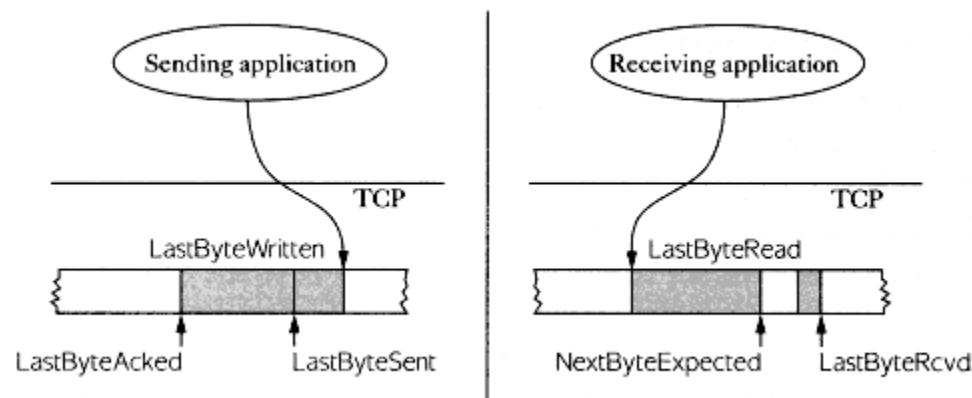


Linux Tx/Rx Ring handling



TCP/IP Details

- The sliding **window** serves several purposes:
- (1) it guarantees the reliable delivery of data
- (2) it ensures that the data is delivered in order,
- (3) it enforces flow control between the sender and the receiver.



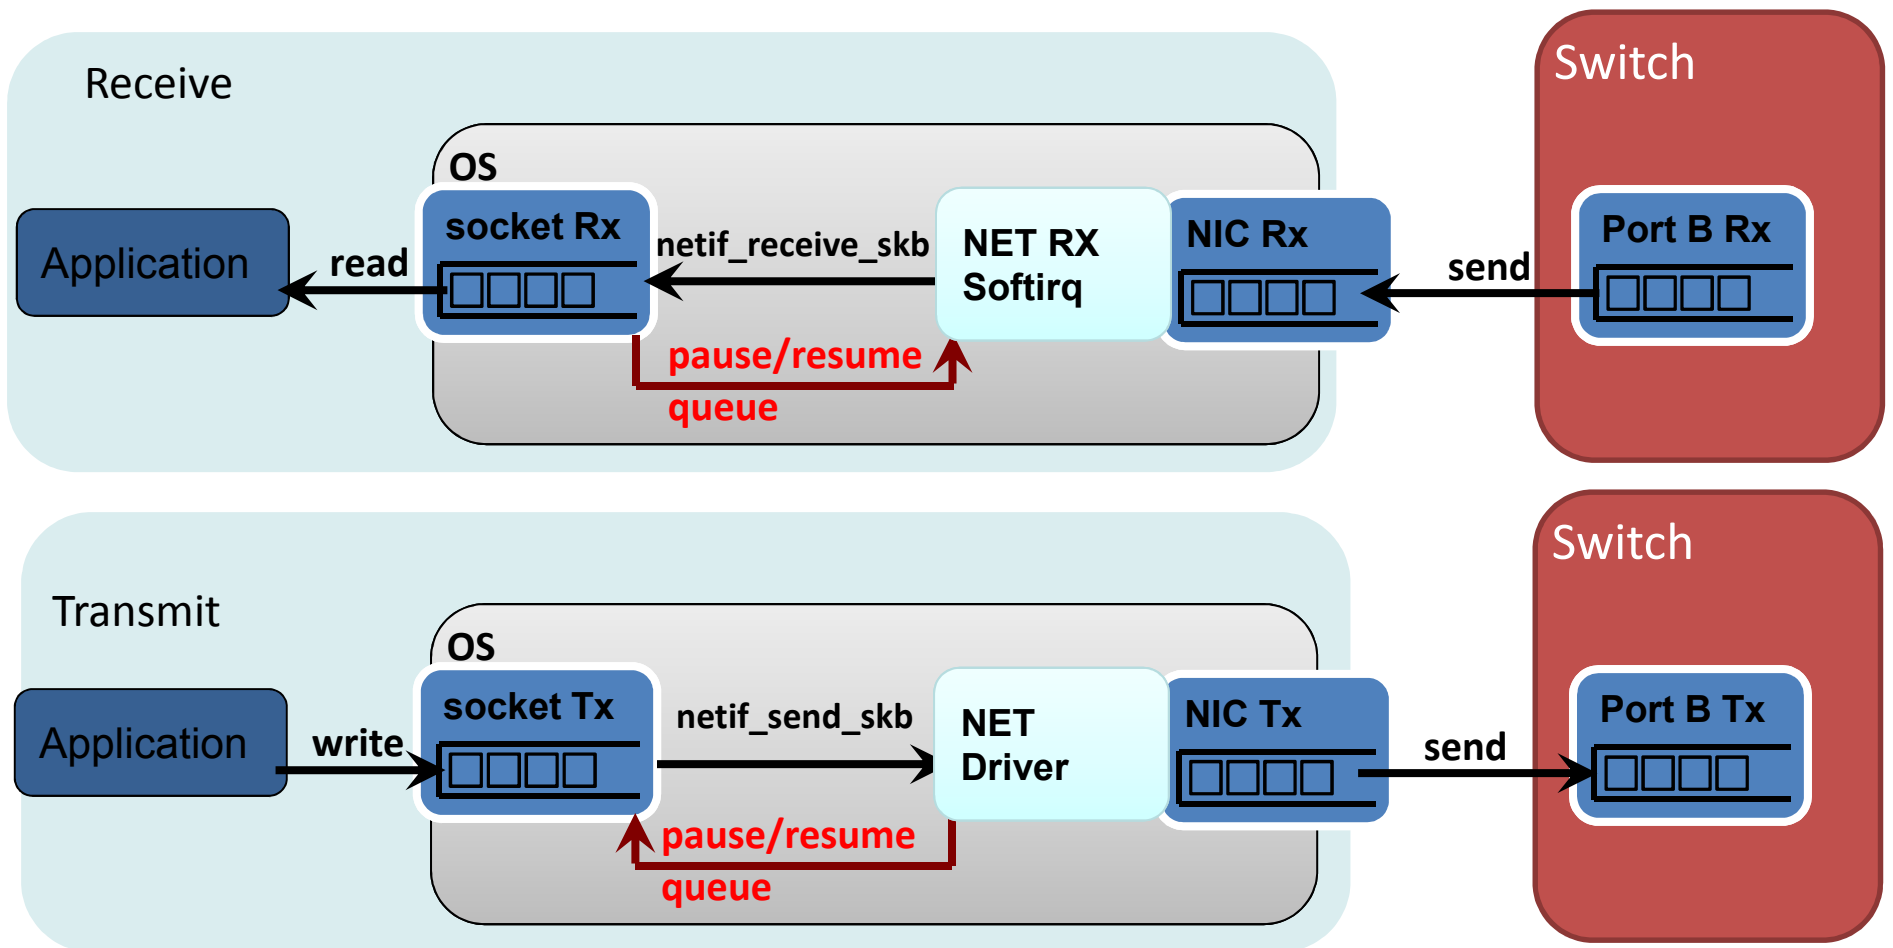
Flow Control

- Max Send and Receive Buffer sizes
- In order delivery to the consumer
- Acknowledgement of reception
- Retransmit when ack is not received in RTT (RoundTripTime) setting

Congestion Control

- Slow Start
 - Start with 1 congestion window and then doubling it
- Fast Retransmit
 - When out of order packet is received immediately ACK
- Fast Recovery

Device Driver Details

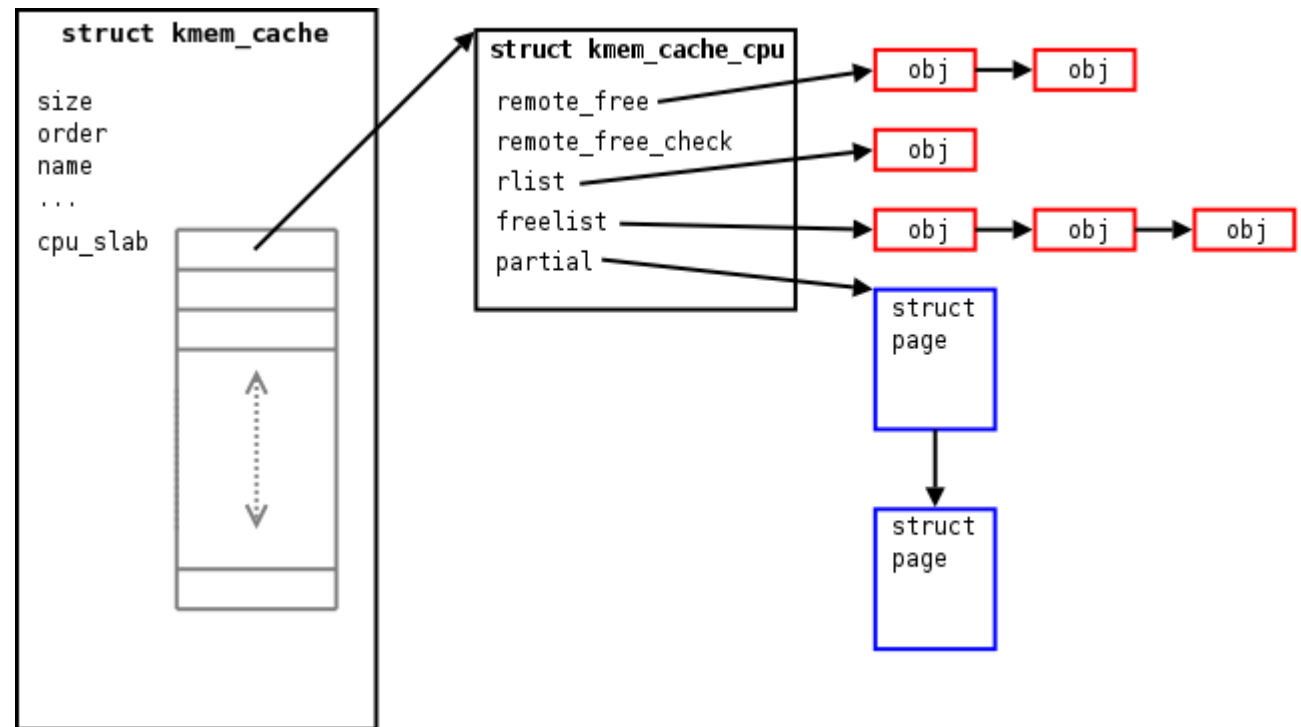


Some useful general "stuff"

- Slab-Cache
 - The primary motivation for slab allocation:
 - initialization and destruction of kernel data objects can actually outweigh the cost of allocating memory for them.
 - As object creation and deletion are widely employed by the kernel, mitigating overhead costs of initialization can result in significant performance gains.
 - "object caching" was therefore introduced in order to avoid the invocation of functions used to initialize object state.
 - Group same dynamically allocated objects under one "allocator" object
 - E.g. `skb_buff_alloc()`

Some useful general "stuff"

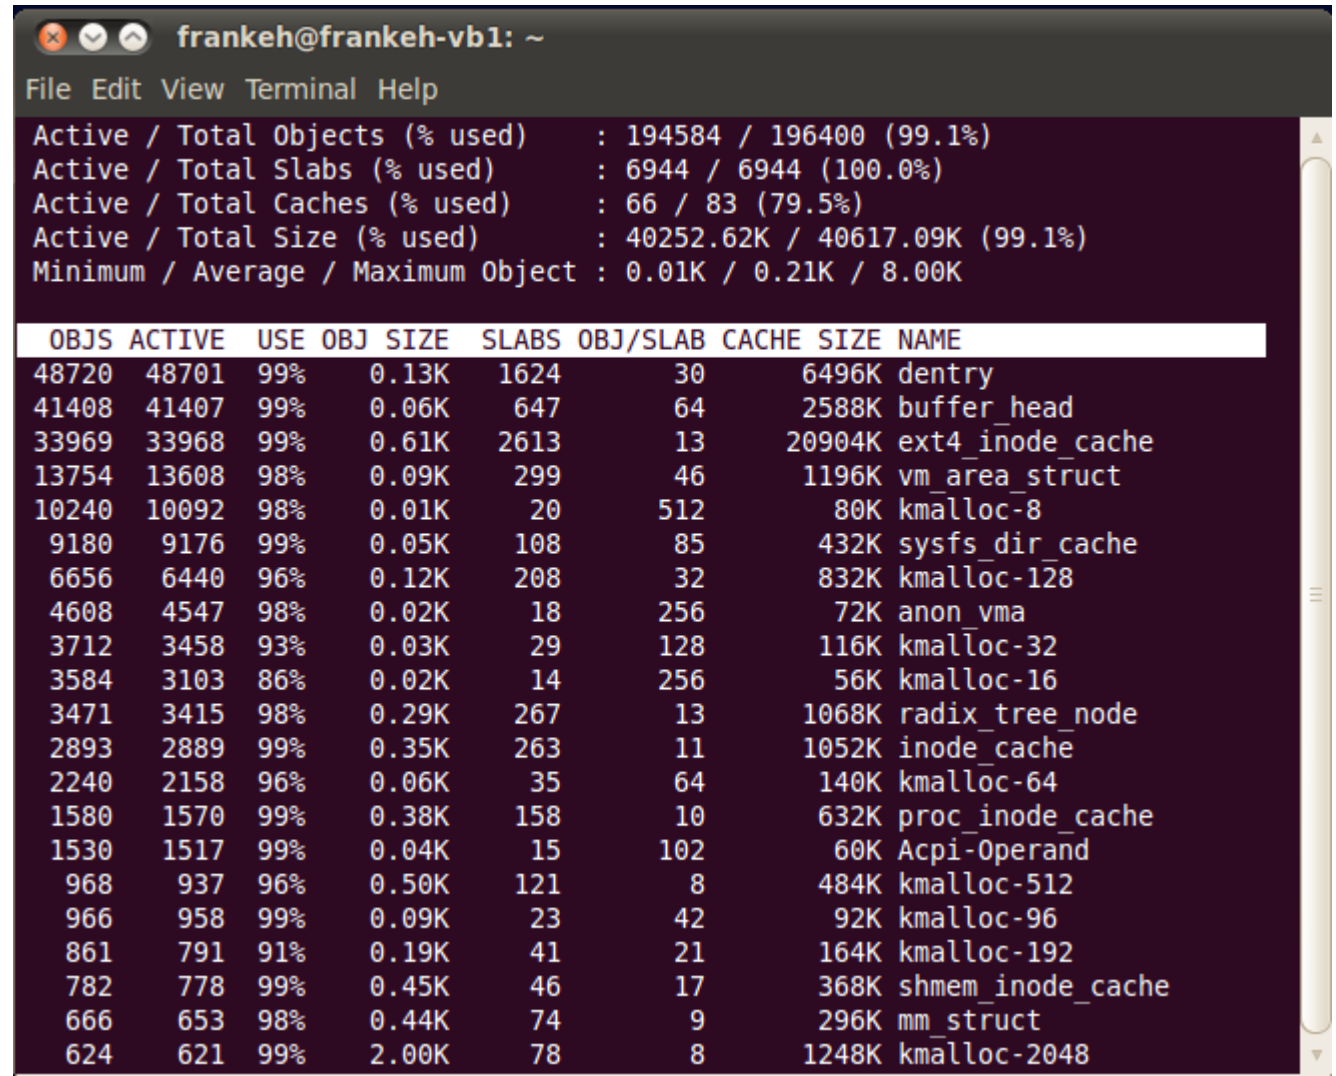
- General implementation



- Other features:
 - Coloring → Cache utilization

Linux Example of slab caches

- #> slabtop



```
frankeh@frankeh-vb1: ~
File Edit View Terminal Help
Active / Total Objects (% used) : 194584 / 196400 (99.1%)
Active / Total Slabs (% used) : 6944 / 6944 (100.0%)
Active / Total Caches (% used) : 66 / 83 (79.5%)
Active / Total Size (% used) : 40252.62K / 40617.09K (99.1%)
Minimum / Average / Maximum Object : 0.01K / 0.21K / 8.00K
```

OBJS	ACTIVE	USE	OBJ	SIZE	SLABS	OBJ/SLAB	CACHE	SIZE	NAME
48720	48701	99%	0.13K	1624	30	6496K	dentry		
41408	41407	99%	0.06K	647	64	2588K	buffer_head		
33969	33968	99%	0.61K	2613	13	20904K	ext4_inode_cache		
13754	13608	98%	0.09K	299	46	1196K	vm_area_struct		
10240	10092	98%	0.01K	20	512	80K	kmalloc-8		
9180	9176	99%	0.05K	108	85	432K	sysfs_dir_cache		
6656	6440	96%	0.12K	208	32	832K	kmalloc-128		
4608	4547	98%	0.02K	18	256	72K	anon_vma		
3712	3458	93%	0.03K	29	128	116K	kmalloc-32		
3584	3103	86%	0.02K	14	256	56K	kmalloc-16		
3471	3415	98%	0.29K	267	13	1068K	radix_tree_node		
2893	2889	99%	0.35K	263	11	1052K	inode_cache		
2240	2158	96%	0.06K	35	64	140K	kmalloc-64		
1580	1570	99%	0.38K	158	10	632K	proc_inode_cache		
1530	1517	99%	0.04K	15	102	60K	Acpi-Operand		
968	937	96%	0.50K	121	8	484K	kmalloc-512		
966	958	99%	0.09K	23	42	92K	kmalloc-96		
861	791	91%	0.19K	41	21	164K	kmalloc-192		
782	778	99%	0.45K	46	17	368K	shmem_inode_cache		
666	653	98%	0.44K	74	9	296K	mm_struct		
624	621	99%	2.00K	78	8	1248K	kmalloc-2048		